

Practical Concurrent and Parallel Programming VIII

Performance Measurements

Jørgen Staunstrup

Performance measurements



Make time consuming computations run faster e.g.:

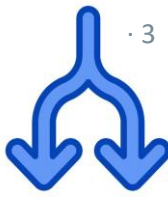
- searching in large amounts of data
- simulation (weather forecasts)
- sorting large volumes of data
- ...

Thread creation is "expensive" (time consuming) !!

```
Thread t= new Thread( ... );
```

Today: we will address how to measure computing times

How long does it take to create a thread?



Thread creation is "expensive" (time consuming) !!

```
Thread t= new Thread( ... );
```

But how expensive?

Assuming that a single floating-point multiplication takes **one** time unit on your PC

**This week we focus on
how to find out.**

**1: (Approximately) How many
time units will it take to create
and start a thread?**



- **Performance measurements: motivation and introduction**
- Pitfalls (and avoiding them)
- Measurements of thread overhead
- Algorithms for parallel computing
- Bonus: Calculating means and variance (efficiently)

Motivations for Concurrency

From Week01



Inherent: User interfaces and other kinds of input/output

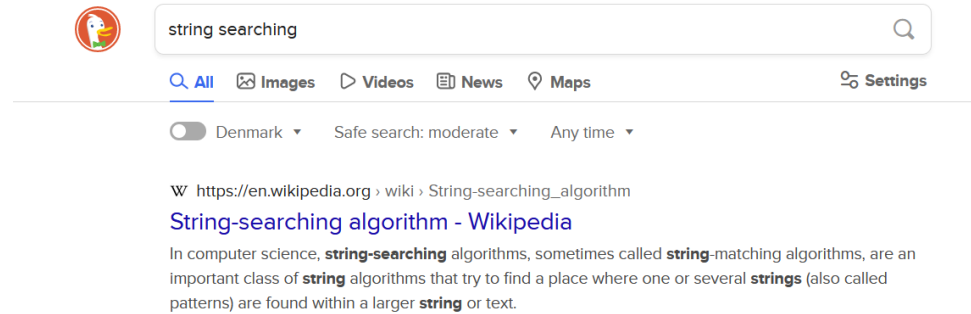
Exploitation: Hardware capable of simultaneously executing multiple streams of statements

Hidden: Enabling several programs to share some resources in a manner where each can act as if they had sole ownership

Motivation 1: Time consuming computations



Searching in a (large) text



<https://www.geeksforgeeks.org/applications-of-string-matching-algorithms/>

Computing prime numbers

2, 3, 5, 7, 11, 13, 17, 19, 23,
29, 31, 37, 41, 43, 47, 53, 59,
61, 67, 71, 73, 79, 83, 89, 97,
...

Cornerstone of all computer security

<https://science.howstuffworks.com/math-concepts/prime-numbers.htm>

Motivation 2: Analyzing code



Thread creation is expensive ?

The Java tutorials say that creating a Thread is expensive. But why exactly is it expensive? What exactly is happening when a Java Thread is created that makes its creation expensive? I'm taking the statement as true, but I'm just interested in mechanics of Thread creation in JVM.

Thread lifecycle overhead. Thread creation and teardown are not free. The actual overhead

But how expensive ?

~ 600 ns to create (on this laptop)

~ 20 times more time than creating a simple object

65000 ns to create and start a thread !!! (on this laptop)

Today: How to get such numbers !

(Performance) Measurements



Key in many sciences (experiments, observations, predictions, ...)

Computer science:

A bit of statistics

A bit of numerical analysis

A bit of computer architecture (cores, caches, number representation,)

Code for measuring execution time

Based on *Microbenchmarks in Java and C#* by Peter Sestoft (see [benchmarkingNotes.pdf](#) in material for this week)

All numbers in these slides were measured in August 2021 on a:

Intel Core i5-1035G4 CPU @ 1.10GHz, 4 Core(s), 8 Logical Processor(s)

- Performance measurements: motivation and introduction
- **Pitfalls (and avoiding them)**
- Measurements of thread overhead
- Algorithms for parallel computing
- Bonus: Calculating means and variance (efficiently)

Example: measuring a (simple) function



```
private static int multiply(int i) {  
    return i * i;  
}
```

```
start= System.nanoTime();  
multiply(126465);  
end= System.nanoTime();  
  
System.out.println(end-start+" ns");
```

Example: measuring a (simple) function



```
start= System.nanoTime();  
multiply(126465);  
end= System.nanoTime();  
  
System.out.println(end-start+" ns");
```

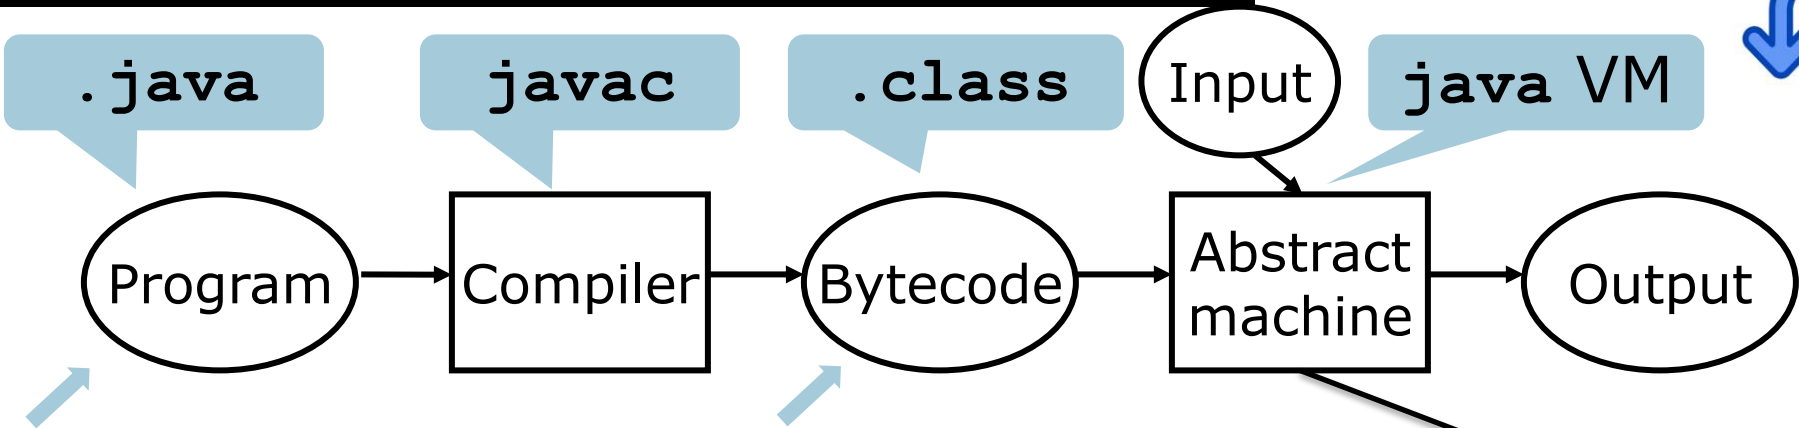
My results: 3600 ns 1400 ns 1500 ns, ...

Should be: ~ 1-2 ns

What is going on?

Java compiler and virtual machine

· 12



```
for (int i=0; i<n; i++)  
    sum += sqrt(arr[i]);
```

```
21 iconst_0  
22 istore 5  
24 iload 5  
26 iload 2  
27 if_icmpge      46  
30 dload 3
```

JIT (Just In Time)

```
19 xorl %ebx,%ebx  
1b jmp 3a  
    0x00(%ebp),%ebp  
    0xec(%ebp)  
    %ebx,0x0c(%edi)  
26 jbe 49  
2c leal  
    0x10(%edi,%ebx,8),%eax  
...
```

Compilation also happens at runtime

```
34 daload  
35 invokestatic Math.sqrt:(D)D  
38 dadd  
39 dstore 3  
40 iinc 5, 1  
43 goto 24
```

JVM

x86

Benchmarking note

· 13



Microbenchmarks in Java and C#

Peter Sestoft (sestoft@itu.dk)

IT University of Copenhagen, Denmark

Version 0.8.0 of 2015-09-16

A goldmine of good advice

Accompanying code: [Benchmark.java](#)

On PCPP GitHub (week08)

Abstract: Sometimes one wants to measure the speed of software, for instance, to measure whether a

```
class Benchmark {
    public static void main(String[] args) { new Benchmark(); }

    public Benchmark() {
        // SystemInfo();
        // Mark0();
        // Mark1();
        ...
        Mark6("multiply", i -> multiply(i));
        ...
        // SortingBenchmarks();
        ...
    }
}
```

Our focus:

Mark functions

The Timer class (in Benchmark.java)



A simple Timer class for Java

Works on all platforms (Linux, MacOS, Windows)

```
public class Timer {  
    private long start, spent = 0;  
    public Timer() { play(); }  
    public double check()  
    { return (System.nanoTime()-start+spent)/1e9; }  
    public void pause() { spent += System.nanoTime()-start; }  
    public void play() { start = System.nanoTime(); }  
}
```

Example: measuring a simple function

· 16



```
private static double multiply(int i) {  
    double x = 1.1 * (double)(i & 0xFF);  
    return x * x * x * x * x * x * x * x * x * x * x * x *  
           * x * x * x * x * x * x * x * x * x * x * x * x;  
}  
  
public static double Mark2() {  
    Timer t = new Timer();  
    int count = 100_000_000;  
    double dummy = 0.0;  
    for (int i=0; i<count; i++)  
        dummy += multiply(i);  
    double time = (t.check() / count) * 1e9 ;  
    System.out.printf("%6.1f ns%n", time);  
    return dummy;  
}
```

Example: measuring a simple function

· 17



```
private static double multiply(int i) {  
    double x = 1.1 * (double)(i & 0xFF);  
    return x * x * x * x * x * x * x * x * x * x * x * x *  
           * x * x * x * x * x * x * x * x * x * x * x * x;  
}  
  
public static double Mark2() {  
    Timer t = new Timer();  
    int count = 100_000_000;  
    double dummy = 0.0;  
    for (int i=0; i<count; i++)  
        dummy += multiply(i);  
    double time = (t.check() / count) * 1e9 ;  
    System.out.printf("%6.1f ns%n", time);  
    return dummy;  
}
```

11.9 ns

Automating multiple runs (Mark3)



Results will usually vary

```
public static double Mark3() {  
    int n = 10;  
    int count = 100_000_000;  
    double dummy = 0.0;  
    for (int j=0; j<n; j++) {  
        Timer t = new Timer();  
        for (int i=0; i<count; i++)  
            dummy += multiply(i);  
        double time = t.check() * 1e9 / count;  
        System.out.printf("%6.1f ns%n", time);  
    }  
    return dummy;  
}
```

```
24.6 ns  
24.6 ns  
24.5 ns  
24.6 ns  
24.4 ns  
24.3 ns  
24.5 ns  
24.4 ns  
24.7 ns  
24.6 ns
```

What is the running time?



What should you report as the result, when the observations are:

30.7 ns 30.3 ns 30.1 ns 30.7 ns 30.5 ns 30.4 ns 30.9 ns 30.3 ns 30.5 ns 30.8 ns ?

Mean: 30.4 ns

What if they are:


30.7 ns 100.2 ns 30.1 ns 30.7 ns 20.2 ns 30.4 ns 2.0 ns 30.3 ns 30.5 ns 5.4 ns ??

Mean: 31.0 ns ??

Normal distribution



· 20



Measuring physical properties

Your exam grades

Course evaluations

Fabrication faults

Running time of Java code

...

Standard deviation/variance



$$\mu = \frac{1}{n} \sum_{j=1}^n t_j$$

$$\sigma = \sqrt{\frac{1}{n-1} \sum_{j=1}^n (t_j - \mu)^2}$$


Mean

Standard deviation


Benchmark note p6

30.7 ns 30.3 ns 30.1 ns 30.7 ns 50.2 ns 30.4 ns 30.9 ns 30.3 ns 30.5 ns 30.8 ns ??

Mean: 32.5 ns Standard deviation: 6.2



Normal distribution



Mark5 - computes mean and variance

· 23




```
public static double Mark5() {
    int n = 10, count = 1, totalCount = 0;
    double dummy = 0.0, runningTime = 0.0, st = 0.0, sst = 0.0;
    do {
        count *= 2;
        st = sst = 0.0;
        for (int j=0; j<n; j++) {
            Timer t = new Timer();
            for (int i=0; i<count; i++) dummy += multiply(i);
            runningTime = t.check();
            double time = runningTime * 1e9 / count;
            st += time;
            sst += time * time;
            totalCount += count;
        }
        double mean = st/n, sdev = Math.sqrt((sst - mean*mean*n) / (n-1));
        System.out.printf("%6.1f ns +/- %8.2f %10d%n", mean, sdev, count);
    } while (runningTime < 0.25 && count < Integer.MAX_VALUE/2);
    return dummy / totalCount;
}
```

Note:
two loops

Mark5 - computes mean and variance



```
public static double Mark5() {
    int n = 10, count = 1, totalCount = 0;
    double dummy = 0.0, runningTime = 0.0, st = 0.0, sst = 0.0;
    do {
        count *= 2;
        st = sst = 0.0;
        for (int j=0; j<n; j++) {
            Timer t = new Timer();
            for (int i=0; i<count; i++) dummy += multiply(i);
            runningTime = t.check();
            double time = runningTime * 1e9 / count;
            st += time;
            sst += time * time;
            totalCount += count;
        }
        double mean = st/n, sdev = Math.sqrt((sst - mean*mean*n)/(n-1));
        System.out.printf("%6.1f ns +/- %8.2f %10d%n", mean, sdev, count);
    } while (runningTime < 0.25 && count < Integer.MAX_VALUE/2);
    return dummy / totalCount;
}
```



```
public static double Mark5() {
    int n = 10, count = 1, totalCount = 0;
    double dummy = 0.0, runningTime = 0.0, st = 0.0, sst = 0.0;
    do {
        count *= 2;
        st = sst = 0.0;
        for (int j=0; j<n; j++) {
            Timer t = new Timer();
            for (int i=0; i<count; i++) dummy += multiply(i);
            runningTime = t.check();
            double time = runningTime * 1e9 / count;
            st += time;
            sst += time * time;
            totalCount += count;
        }
        double mean = st/n, sdev = Math.sqrt((sst - mean*mean*n)/(n-1));
        System.out.printf("%6.1f ns +/- %8.2f %10d%n", mean, sdev, count);
    } while (runningTime < 0.25 && count < Integer.MAX_VALUE/2);
    return dummy / totalCount;
}
```


Parameterizing function to be measured



```
private static double multiply(int i) {  
    . . .  
}
```

Java: `multiply(i)` is a number

Java: `i -> multiply(i)` is a function

<https://docs.oracle.com/javase/tutorial/java/javaOO/lambdaexpressions.html>

```
Mark6( . . . , i -> multiply(i));
```

Mark6 - introduce a functional argument



· 27

The function **f** is benchmarked

```
public static double Mark6(String msg, IntToDoubleFunction f) {
    int n = 10, count = 1, totalCount = 0;
    double dummy = 0.0, runningTime = 0.0, st = 0.0, sst = 0.0;
    do {
        count *= 2;
        st = sst = 0.0;
        for (int j=0; j<n; j++) {
            Timer t = new Timer();
            for (int i=0; i<count; i++) dummy += f.applyAsDouble(i);
            runningTime = t.check();
            double time = runningTime * 1e9 / count;
            st += time; sst += time * time; totalCount += count;
        }
        double mean = st/n, sdev = Math.sqrt((sst - mean*mean*n)/(n-1));
        System.out.printf("%-25s %15.1f ns %10.2f %10d%n", msg, mean, sdev, count);
    } while (runningTime < 0.25 && count < Integer.MAX_VALUE/2);
    return dummy / totalCount;
}

public interface IntToDoubleFunction { double applyAsDouble(int i); }

Mark6("multiply", i -> multiply(i));
```

←
lambda

Example use of Mark6

```
Mark6("multiply", i -> multiply(i));
```

| | | | |
|----------|----------|---------|----------|
| multiply | 595.0 ns | 1407.81 | 2 |
| multiply | 147.5 ns | 90.10 | 4 |
| multiply | 212.5 ns | 152.53 | 8 |
| multiply | 170.6 ns | 59.44 | 16 |
| multiply | 201.9 ns | 157.69 | 32 |
| multiply | 60.8 ns | 34.55 | 64 |
| multiply | 65.1 ns | 59.83 | 128 |
| multiply | 54.3 ns | 14.85 | 256 |
| ... | | | |
| multiply | 24.6 ns | 0.75 | 524288 |
| multiply | 24.6 ns | 0.88 | 1048576 |
| multiply | 24.9 ns | 2.71 | 2097152 |
| multiply | 24.3 ns | 0.85 | 4194304 |
| multiply | 24.2 ns | 0.72 | 8388608 |
| multiply | 25.0 ns | 1.38 | 16777216 |

iterations



Mark7 - printing only final values



```
public static double Mark7(String msg, IntToDoubleFunction f) {  
    ...  
    do {  
        ...  
    } while (runningTime < 0.25 && count < Integer.MAX_VALUE/2);  
    double mean = st/n, sdev = Math.sqrt((sst - mean*mean*n)/(n-1));  
    System.out.printf("%-25s %15.1f %10.2f %10d%n", msg, mean, sdev, count);  
    return dummy / totalCount;  
}
```

- Performance measurements: motivation and introduction
- Pitfalls (and avoiding them)
- **Measurements of thread overhead**
- Algorithms for parallel computing
- Bonus: Calculating means and variance (efficiently)



```
Mark7("Thread create",
    i -> {
        Thread t= new Thread(() -> {
            for (int j= 0; j<1000; j++)
                ai.getAndIncrement();
        });
        return t.hashCode(); // to confuse compiler to not optimize
    });
```

Takes 620 ns

2: What are we really measuring?

Thread creation

· 32



```
Mark7("Thread create",
    i -> {
        Thread t= new Thread(() -> {
            for (int j= 0; j<1000; j++) // not executed
                ai.getAndIncrement(); // thread t created, but not started
        });
        return t.hashCode(); // to confuse compiler to not optimize
    });
```

Takes 620 ns

Creating a thread/an object +
return statement



```
Mark7("Thread create",
  i -> {
    Thread t= new Thread(() -> {
      for (int j= 0; j<1000; j++) // not executed
        ai.getAndIncrement();    // thread t created, but not started
    });
    return t.hashCode(); // to confuse compiler to not optimize
  });
```

Find out running times for:

1. creating an object
2. calculating hash-code
3. starting a thread

Creating an object

· 34



A thread is an object, so let us start finding the cost of creating a simple object.

```
class Point {  
    public final int x, y;  
    public Point(int x, int y) { this.x = x; this.y = y; }  
}
```

```
Mark7("hashCode()", i -> myPoint.hashCode());
```

```
Mark7("Point creation",  
    i -> {  
        Point p= new Point(i, i);  
        return p.hashCode();  
    });
```

hashCode() 3 ns

Point creation 33 ns

So, object creation is: ~ 30 ns

Thread create + start



```
Mark7(" ... ",  
    i -> {  
        Thread t= new Thread(() -> {  
            for (int j= 0; j<1000; j++)  
                ai.getAndIncrement();  
        });  
        t.start();  
        return t.hashCode();  
    });
```

3: What are we really measuring?

Thread create + start



```
Mark7(" ... ",  
    i -> {  
        Thread t= new Thread(() -> {  
            for (int j= 0; j<1000; j++)  
                ai.getAndIncrement();  
        });  
        t.start();  
        return t.hashCode();  
    });
```

For loop not included!!!

Thread create + start



```
Mark7(" ... ",  
    i -> {  
        Thread t= new Thread(() -> {  
            for (int j= 0; j<1000; j++)  
                ai.getAndIncrement();  
        });  
        t.start();  
        return t.hashCode();  
    });
```

lambda returns right after start

Thread create + start

· 38



```
Mark7(" ... ",  
      i -> {  
          Thread t= new Thread(() -> {  
              for (int j= 0; j<1000; j++)    //most iterations not done  
                  ai.getAndIncrement();      // Why?  
          });  
          t.start();  
          return t.hashCode();  
      });
```

Takes ~ 65000 ns

- So, a lot of work goes into **starting** a thread
- Even after creating it (takes ~ 600 ns)
- Note: does not include executing the for loop

Never create threads for small computations !!!

- Performance measurements: motivation and introduction
- Pitfalls (and avoiding them)
- Measurements of thread overhead
- **Algorithms for parallel computing**
- Bonus: Calculating means and variance (efficiently)



Quicksort: <https://www.chrislaux.com/quicksort.html>

```
private static void qsort(int[] arr, int a, int b) {
    if (a < b) {
        int i = a, j = b;
        int x = arr[(i+j) / 2];
        do {
            while (arr[i] < x) i++;
            while (arr[j] > x) j--;
            if (i <= j) { swap(arr, i, j); i++; j--; }
        } while (i <= j);
        qsort(arr, a, j); qsort(arr, i, b);
    }
}
```

see SearchAndSort.java in week08 material

Prime counting: <https://www.dcode.fr/prime-number-pi-count>

```
long count = 0;
final int from = 0, to = range;
for (int i=from; i<to; i++) if (isPrime(i)) count++;
```

Multithreaded version of CountPrimes



· 41

2, 3, 4, 5,



thread0



thread1

range



threadN

Code for exercises week08: [TestCountPrimesThreads.java](#)

Java code for TestCountTimesThreads

42



```
private static long countParallelN(int range, int threadCount) {
    final int perThread= range / threadCount;
    final LongCounter lc= new PrimeCounter();
    Thread[] threads= new Thread[threadCount];
    for (int t=0; t<threadCount; t++) {
        final int from= perThread * t,
            to= (t+1==threadCount) ? range : perThread * (t+1);
        threads[t]= new Thread()->
            {for (int i=from; i<to; i++)
                if (isPrime(i)) lc.increment();
            });
    }
    for (int t=0; t<threadCount; t++) threads[t].start();
    try { for (int t=0; t<threadCount; t++) threads[t].join();
    } catch (InterruptedException exn) { }
    return lc.get();
}
```

Running times for prime counting



| | | | |
|---------------|---|--------------|-----------|
| countParallel | 1 | 7107236.6 ns | 448417.55 |
| countParallel | 2 | 6069944.7 ns | 802224.61 |
| countParallel | 3 | 3621185.5 ns | 152693.03 |
| countParallel | 4 | 3124067.0 ns | 640480.51 |

...

Good or bad ?

**Number of primes are very different
in these ranges !!!**

2, 3, 4, 5,



thread0



thread1



range

threadN

Breaking the task into smaller pieces/tasks



2, 3, 4, 5,



range



When a thread is done with one task, it gets a new task
until all tasks are done



[low..high]

```
when high-low > threshold  
int mid= low+(high-low)/2;
```

else sequential count

[a..mid]
[mid+1..high]

Prime counter task (skeleton)

47



```
public class countPrimesTask implements Runnable {  
    private final int low;  
    private final int high;  
    private final ExecutorService pool;  
  
    @Override public void run() {  
  
        int mid= low+(high-low)/2;  
        pool.submit( new countPrimesTask(low, mid, pool) );  
        pool.submit( new countPrimesTask(mid+1, high, pool) );  
  
    }  
}
```

More on executors next week

Shortcomings:

1. How to stop?
2. Will create too many "small" tasks
3. Returning result (# primes)




- Performance measurements: motivation and introduction
- Pitfalls (and avoiding them)
- Measurements of thread overhead
- Algorithms for parallel computing
- **Bonus: Calculating means and variance (efficiently)**

Normal distribution



· 49



Measuring physical properties

Your exam grades


Course evaluations

Fabrication faults

Running time of Java code


...

But there are exceptions



Source: Bent Flyvbjerg, Alexander Budzier, Jong Seok Lee, Mark Keil, Daniel Lunn & Dirk W. Bester (2022) The Empirical Reality of IT Project Cost Overruns: Discovering A Power-Law Distribution, Journal of Management Information Systems, 39:3, 607-639, DOI: 10.1080/07421222.2022.2096544

Normal distribution



What should you report as the result, when the observations are:

30.7 ns 30.3 ns 30.1 ns 30.7 ns 50.2 ns 30.4 ns 30.9 ns 30.3 ns 30.5 ns 30.8 ns ??

Mean: 32.5 ns Standard deviation: 6.2

50.2 is an outlier

because there is a probability of less than 4.6 % that 50.2 is a valid observation

Computing the variance



$$\mu = \frac{1}{n} \sum_{j=1}^n t_j$$

$$\sigma = \sqrt{\frac{1}{n-1} \sum_{j=1}^n (t_j - \mu)^2}$$

Requires two passes through the data

$$\sigma^2 = \frac{1}{n(n-1)} (n \sum_{j=1}^n t_j^2 - (\sum_{j=1}^n t_j)^2)$$

Can be done in one pass (on-line alg.)

```
for (int j=0; j<n; j++) {  
    Timer t = new Timer();  
    for (int i=0; i<count; i++)  
        ...  
    double time = t.check() * 1e9 / count;  
    st += time;  
    sst += time * time;  
}  
double mean = st/n, sdev = Math.sqrt((sst - mean*mean*n) / (n-1));  
System.out.printf("%6.1f ns +/- %6.3f%n", mean, sdev);
```

The two formulas give the same result



$$\mu = \frac{1}{n} \sum_{j=1}^n t_j$$

$$\sigma = \sqrt{\frac{1}{n-1} \sum_{j=1}^n (t_j - \mu)^2}$$

$$\sigma = \sqrt{\frac{1}{n-1} \sum_{j=1}^n (t_j^2 + \mu^2 - 2t_j\mu)}$$

$$\sigma^2 = \frac{1}{n-1} \sum_{j=1}^n (t_j^2 + \mu^2 - 2t_j\mu)$$

$$\sigma^2 = \frac{1}{n-1} (\sum_{j=1}^n t_j^2 + \sum_{j=1}^n (\mu^2 - 2t_j\mu))$$

$$\sigma^2 = \frac{1}{n-1} (\sum_{j=1}^n t_j^2 + n\mu^2 - 2\mu \sum_{j=1}^n t_j)$$

$$\sigma^2 = \frac{1}{n-1} (\sum_{j=1}^n t_j^2 + n\mu^2 - 2\mu n\mu)$$

$$\sigma^2 = \frac{1}{n-1} (\sum_{j=1}^n t_j^2 - n\mu^2)$$

$$\sigma^2 = \frac{1}{n(n-1)} (n \sum_{j=1}^n t_j^2 - \mu^2)$$

$$\sigma^2 = \frac{1}{n(n-1)} (n \sum_{j=1}^n t_j^2 - (\frac{1}{n} \sum_{j=1}^n t_j)^2)$$



Formula in Benchmark note

See exercises08.pdf

also https://en.wikipedia.org/wiki/Algorithms_for_calculating_variance



Formula used in code (one pass algorithm)



$$\sigma^2 = \frac{1}{n(n-1)} \left(n \sum_{i=1}^n x_i^2 - \left(\sum_{i=1}^n x_i \right)^2 \right)$$

```
int n = 10;
...
for (int j=0; j<n; j++) {
    Timer t = new Timer();
    for (int i=0; i<count; i++)
        ...
    double time = t.check() * 1e9 / count;
    st += time;
    sst += time * time;
}
double mean = st/n, sdev = Math.sqrt((sst - mean*mean*n) / (n-1));
System.out.printf("%6.1f ns +/- %6.3f%n", mean, sdev);
```

Beware: $\text{sst} - \text{mean} * \text{mean} * n$

can be a very small number

Beware of cancellation when subtracting numbers that are close to each other:

28 significant digits

$$\begin{array}{r} 1010101000010110110001110101.111 \\ -1010101000010110110001110001.100 \\ \hline 0000000000000000000000000100.011 \end{array}$$

3 significant digits

<https://blog.demofox.org/2017/11/21/floating-point-precision>